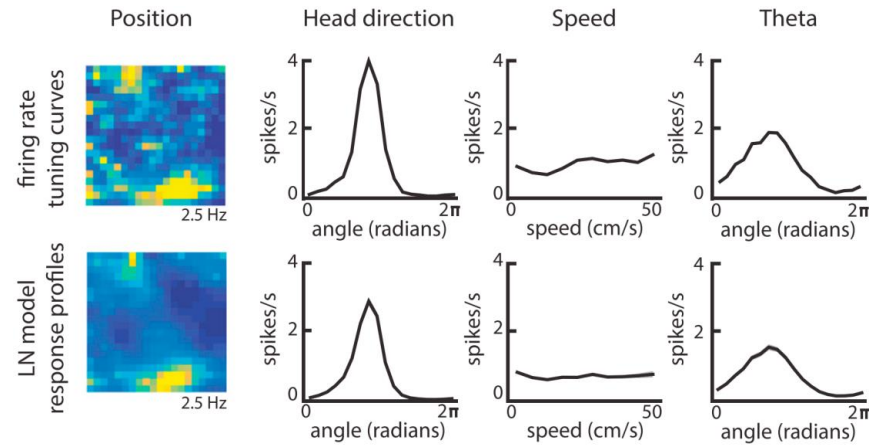


# Generalized linear models



$$\mathbf{W}\mathbf{x} = \mathbf{W}^P\mathbf{x}_P + \mathbf{W}^H\mathbf{x}_H + \mathbf{W}^S\mathbf{x}_S + \mathbf{W}^T\mathbf{x}_T$$

Machine Learning from Scratch seminar

John Vastola

2/20/24

## **Q.** What is **Machine Learning from Scratch**?

**A.** Occasional HMS seminar focused on learning about ML fundamentals.  
Topics might be related to neuro research, or might just be cool!

<https://github.com/DrugowitschLab/ML-from-scratch-seminar>

Seminars happen over 2 evenings:

first is **theory-focused**  
second is **coding-focused**

~ **Four** seminars planned for Spring 2024:

Feb | **GLMs** (w/ Kiah Hardcastle)

Mar | **RNNs** (w/ Siyan Zhou)

Apr | **TBD**

May | **TBD**

# Useful resources for learning about GLMs

## Tutorials

**2016 SFN tutorial on GLMs** by Jesse Kaminsky and Jonathan Pillow  
[https://github.com/pillowlab/GLMspiketraintutorial\\_python](https://github.com/pillowlab/GLMspiketraintutorial_python)

**Neuromatch Academy GLM tutorial** by Fiquet et al.  
[https://compneuro.neuromatch.io/tutorials/W1D3\\_GeneralizedLinearModels/student/W1D3\\_Tutorial1.html](https://compneuro.neuromatch.io/tutorials/W1D3_GeneralizedLinearModels/student/W1D3_Tutorial1.html)

**GLM\_Tensorflow2 repository** by Shih-Yi Tseng  
[https://github.com/sytseng/GLM\\_Tensorflow\\_2](https://github.com/sytseng/GLM_Tensorflow_2)

See this session's GitHub page for more!

# Useful resources for learning about GLMs

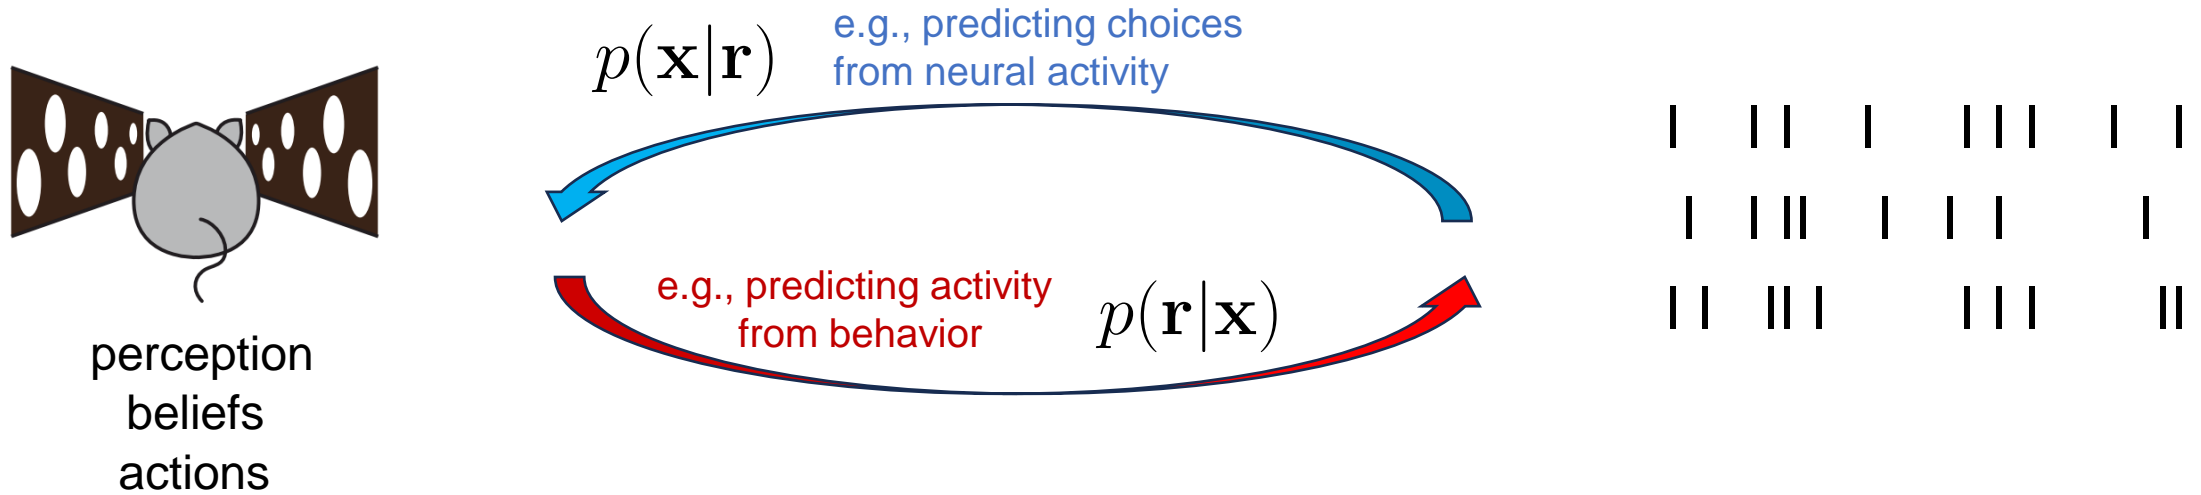
## Example papers

[2014 Park et al.] “Encoding and decoding in parietal cortex during sensorimotor decision-making”  
application to data from macaque lateral intraparietal area  
<https://www.nature.com/articles/nn.3800>


[2017 Hardcastle et al.] “A Multiplexed, Heterogeneous, and Adaptive Code for Navigation in Medial Entorhinal Cortex”  
application to data from mouse medial entorhinal cortex  
<https://doi.org/10.1016/j.neuron.2017.03.025>

[2022 Tseng and Chettih et al.] “Shared and specialized coding across posterior cortical areas for dynamic navigation decisions”  
application to data from mouse posterior cortex  
<https://doi.org/10.1016/j.neuron.2022.05.012>


See this session’s GitHub page for more!




**Q. How do we estimate  $p(\mathbf{r} | \mathbf{x})$  using a mix of neural and behavioral data?**



Lines are simple + interpretable...



...but not necessarily a good fit for neural data!

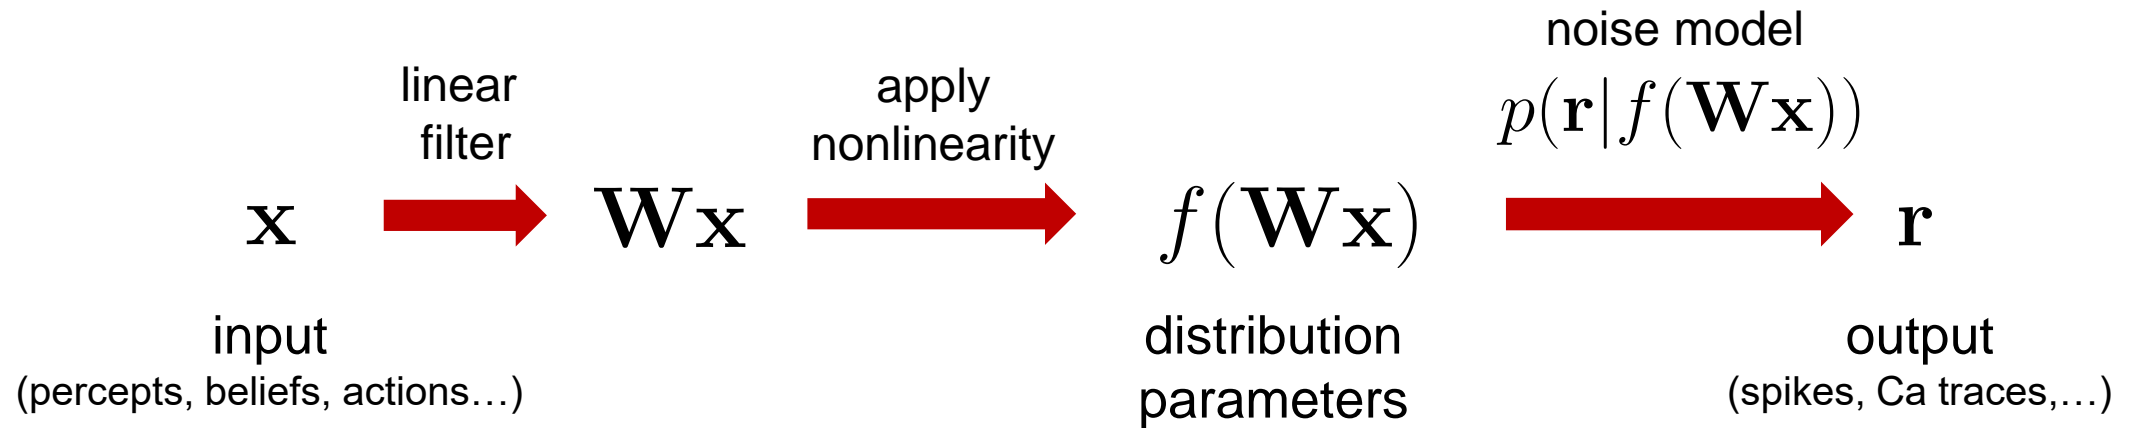


## Solution

Need to add nonlinearity to account for, e.g., nonnegativity.

Also need to allow for non-Gaussian (in particular, Poisson-like) noise.

# GLMs from a generative modeling perspective



## Examples:

- **linear regression**

$$f(\mathbf{z}) = \mathbf{z}$$

$$p(\mathbf{r} | \boldsymbol{\mu}) = \mathcal{N}(\boldsymbol{\mu}, \sigma^2 \mathbf{I})$$

- **logistic regression**

$$f(\mathbf{z}) = 1 / (1 + e^{-\mathbf{z}})$$

$$p(r | p) = p^r (1 - p)^{1-r}$$

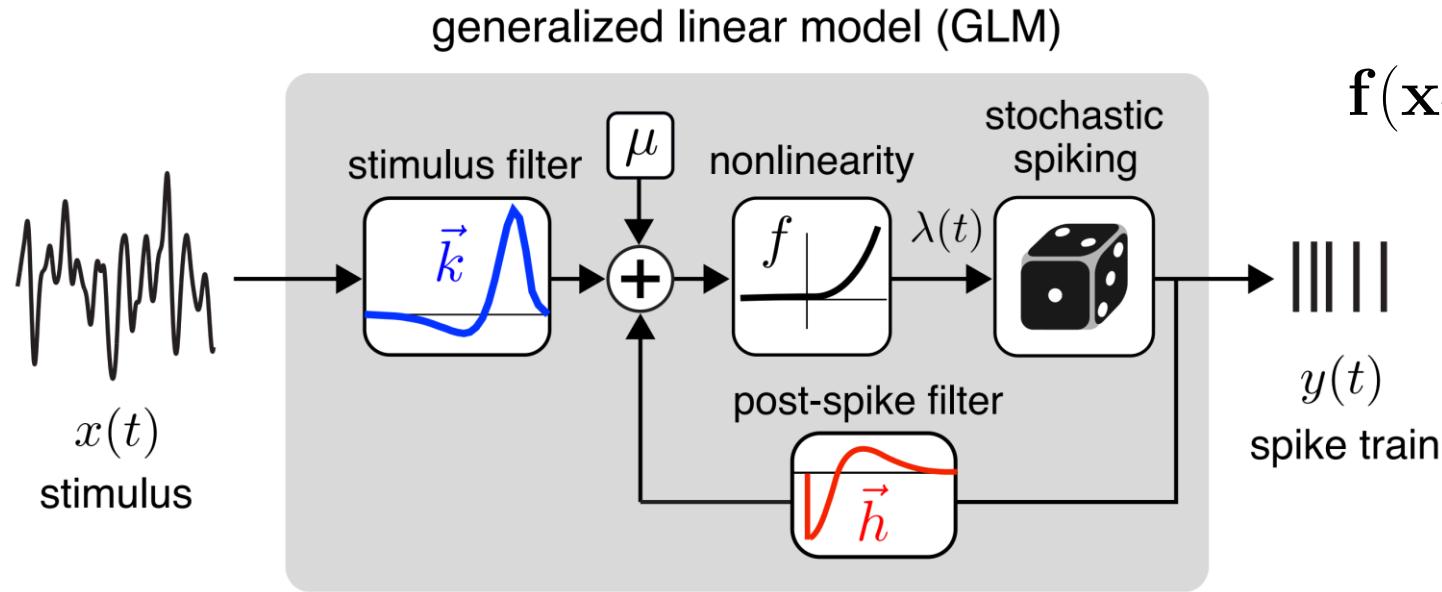
- **Poisson regression**

$$f(\mathbf{z}) = \exp(\mathbf{z})$$

$$p(\mathbf{r} | \boldsymbol{\mu}) = \mathcal{P}(\boldsymbol{\mu})$$

**Note for math lovers:** These approaches can be unified using the idea of exponential family distributions. See Murphy, "Probabilistic Machine Learning" Book 1, Ch. 12. <https://probml.github.io/pml-book/>

# Neuroscience GLMs in practice: linear-nonlinear-Poisson models



$$\mathbf{f}(\mathbf{x}_t, \{\mathbf{y}_{t'}\}) = \exp\left[ \mathbf{W}\mathbf{x}_t + \sum_{t' < t} \mathbf{W}^{t'} \mathbf{y}_{t'} \right]$$

$$\mathbf{y}_t \sim \mathcal{P}(\mathbf{f}(\mathbf{x}_t, \{\mathbf{y}_{t'}\})\Delta t)$$

**Poisson regression** (i.e., using an independent Poisson noise model) most commonly used GLM.

The canonical choice of  $\mathbf{f}()$  is exponential (see exp. family math) but others sometimes used.

Using recent spiking history as input is optional, but helps capture, e.g., bursting, refractory period...



# Poisson regression details: likelihood and gradients

## Log-likelihood

$$\log p = \sum_{n,i} r_i^{(n)} \sum_j W_{ij} x_j^{(n)} - f(\sum_k W_{ik} x_k^{(n)}) \Delta t$$

n: samples, i: neurons, j: inputs

## Gradient wrt weights

$$\begin{aligned} \frac{\partial \log p}{\partial W_{ij}} &= \sum_n \left[ r_i^{(n)} - f(\sum_k W_{ik} x_k^{(n)}) \Delta t \right] x_j^{(n)} \\ &= \sum_n \left[ r_i^{(n)} - \hat{r}_i^{(n)} \right] x_j^{(n)} \end{aligned}$$

Note simplicity of expression:

the nth sample contributes a term (true – estimated)\*input. Two-factor learning rule!

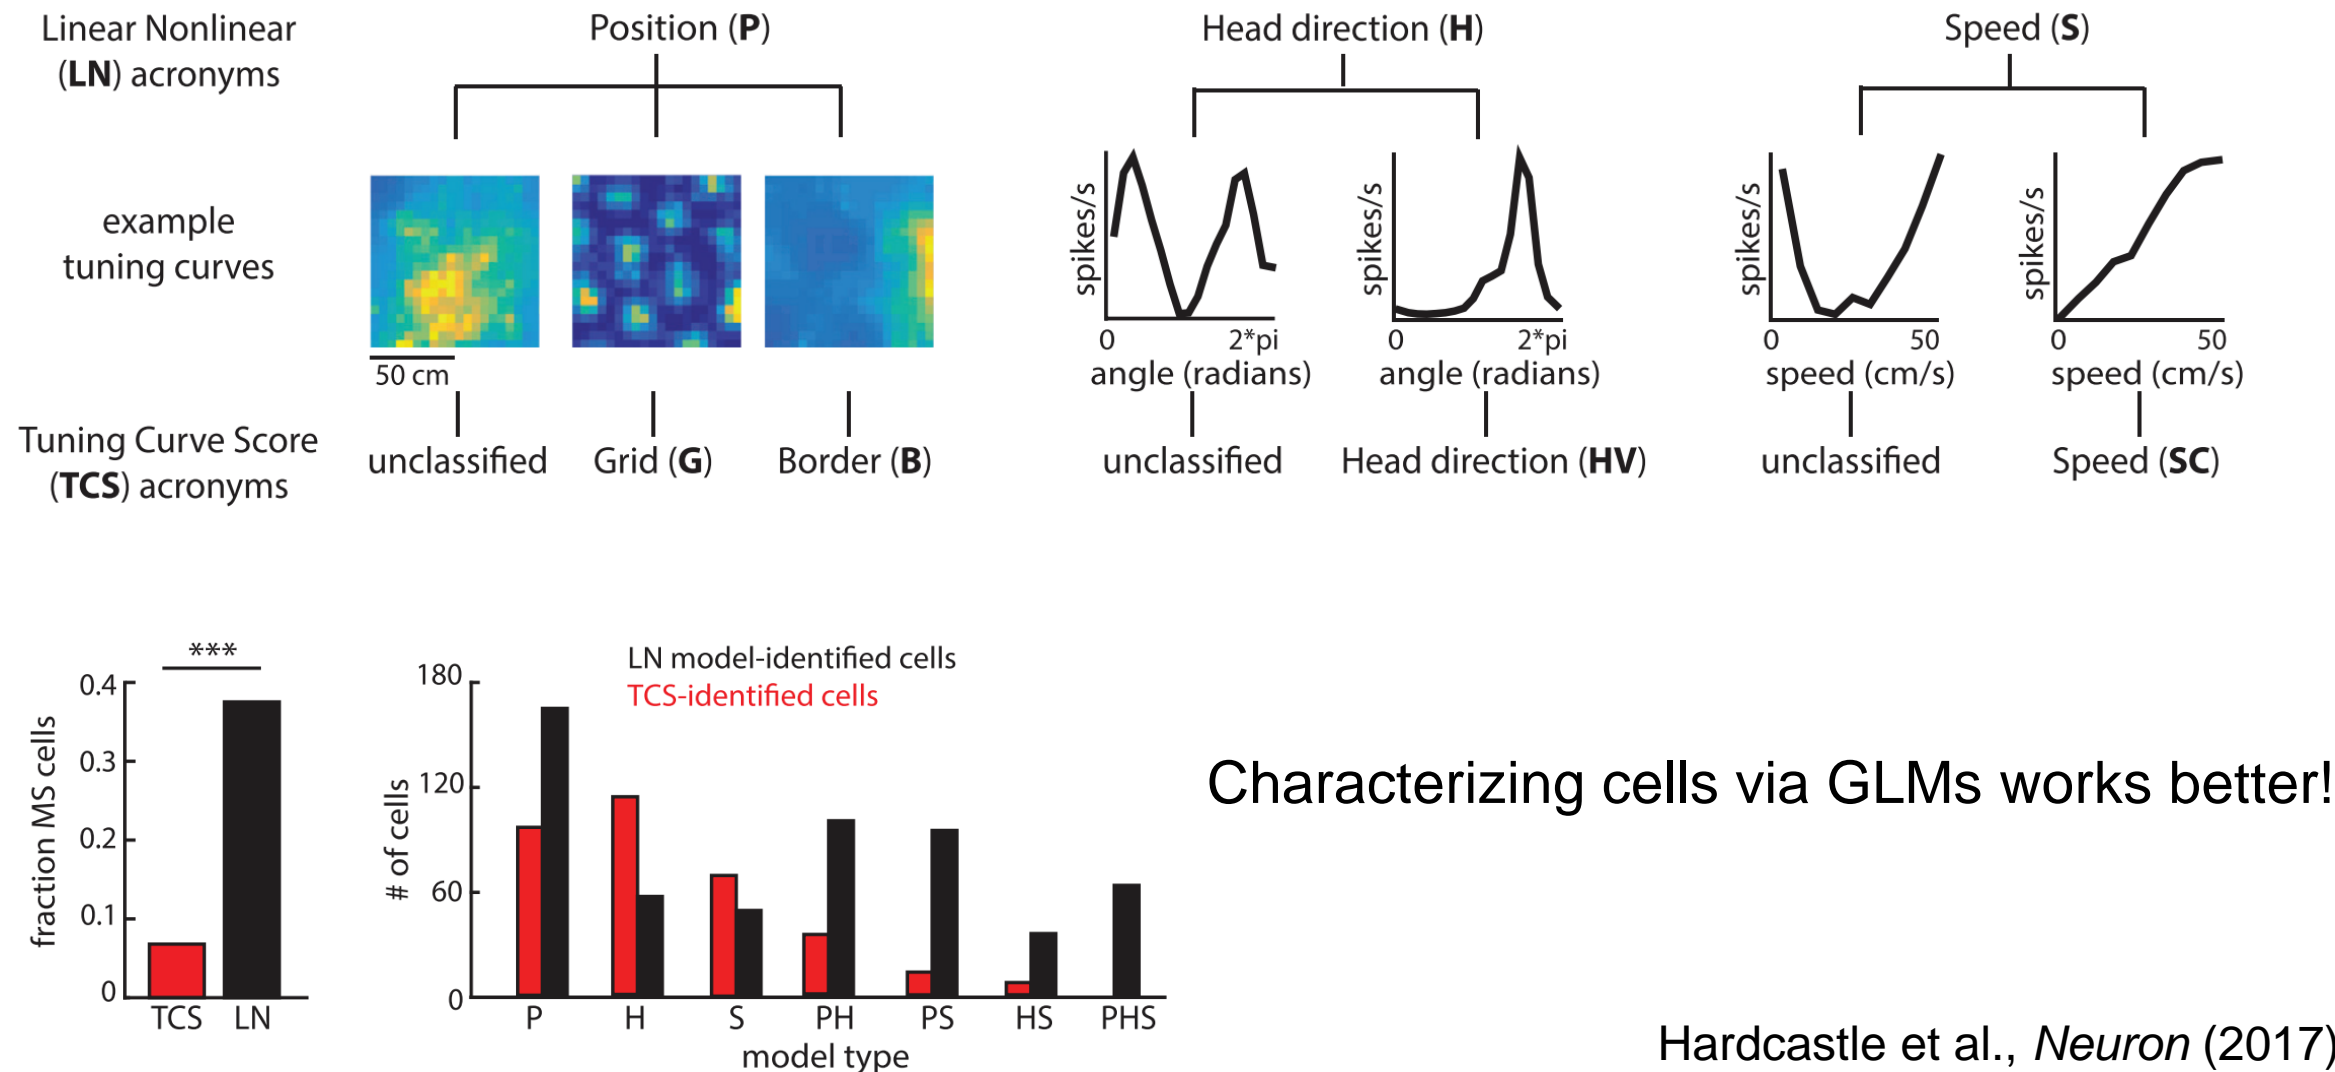
This holds for more general GLMs with canonical activation functions too. See Murphy.

## Example 1: Fitting data from mouse medial entorhinal cortex

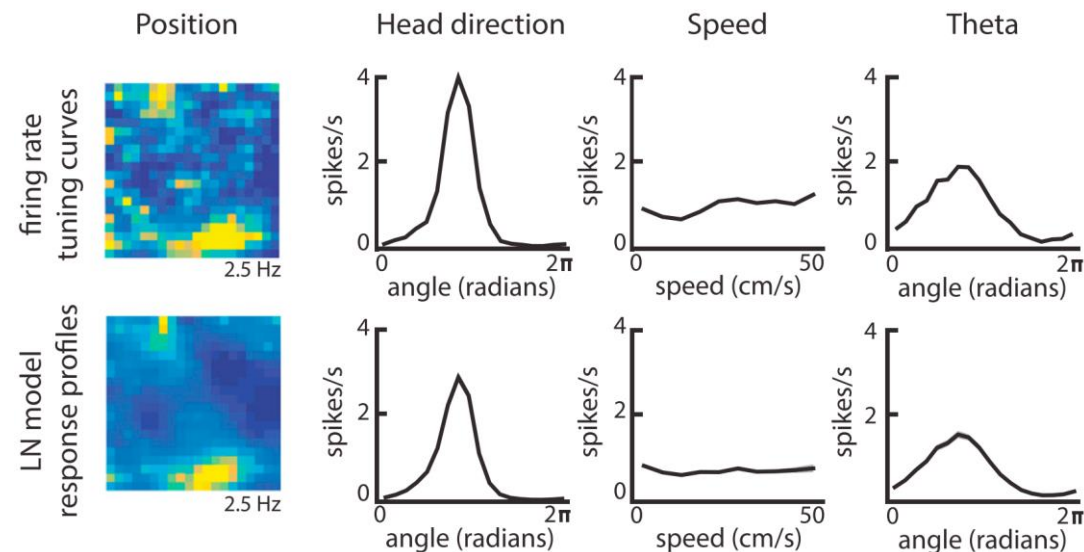
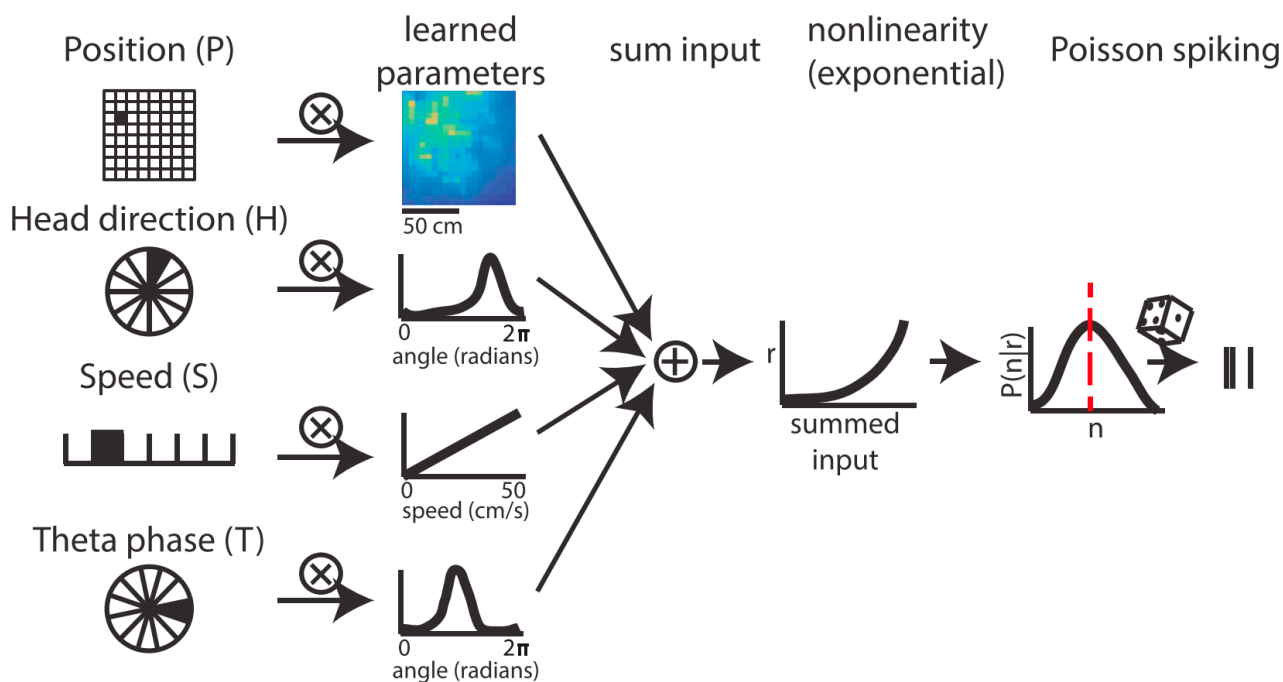
Hardcastle et al., *Neuron* (2017)

“A Multiplexed, Heterogeneous, and Adaptive Code  
for Navigation in Medial Entorhinal Cortex”

# Example 1: Fitting data from mouse medial entorhinal cortex



# Example 1: Fitting data from mouse medial entorhinal cortex



$$p(\mathbf{r}|\mathbf{x}) = \mathcal{P}(e^{\mathbf{W}\mathbf{x}}\Delta t)$$

Poisson observations,  
exponential activation function

$$\mathbf{W}\mathbf{x} = \mathbf{W}^P\mathbf{x}_P + \mathbf{W}^H\mathbf{x}_H + \mathbf{W}^S\mathbf{x}_S + \mathbf{W}^T\mathbf{x}_T$$

position


head-direction

speed

theta phase

# Example 1: Fitting data from mouse medial entorhinal cortex

Position (P)



Head direction (H)



Speed (S)



Theta phase (T)



**Important fitting detail:**

Since behavior encodings are **'one-hot'**, need to regularize objective with smoothness penalty that ensures firing rates for nearby bins are related:

$$\begin{aligned} & \beta_P \sum_{i,j} \frac{(W_{ij}^P - W_{i,j+1}^P)^2}{2} + \beta_H \sum_{i,j} \frac{(W_{ij}^H - W_{i,j+1}^H)^2}{2} \\ & + \beta_S \sum_{i,j} \frac{(W_{ij}^S - W_{i,j+1}^S)^2}{2} + \beta_T \sum_{i,j} \frac{(W_{ij}^T - W_{i,j+1}^T)^2}{2} \end{aligned}$$

Beta coefficients control strength of smoothness penalty for each variable.

# Example 1 Takeaways

- GLMs provide a **statistically unbiased**\* way to characterize tuning of neurons. Better way to characterize MEC cells than with grid scores, border scores, etc.

Caveat: depends somewhat on state representation used as input.

Hardcastle et al. use maximally expressive one-hot behavior encodings to avoid issues.

- Regularization is important for avoiding artifacts since data is limited. Without it, would tend to get weird/discontinuous tuning maps.

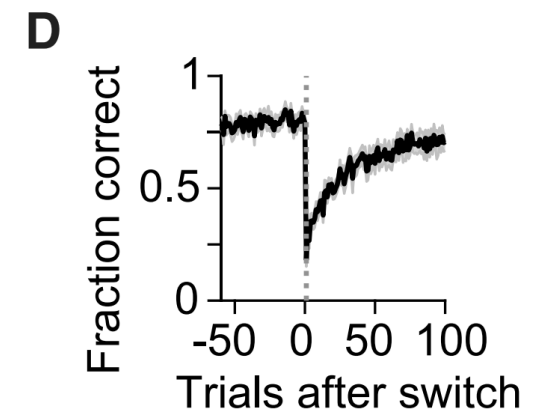



Caveat: as we will see, can also address smoothness using a different state representation.

## Example 2: Fitting data from mouse posterior cortex

Tseng and Chettih et al., *Neuron* (2022)

“Shared and specialized coding across posterior cortical areas for dynamic navigation decisions”

## Example 2: Fitting data from mouse posterior cortex



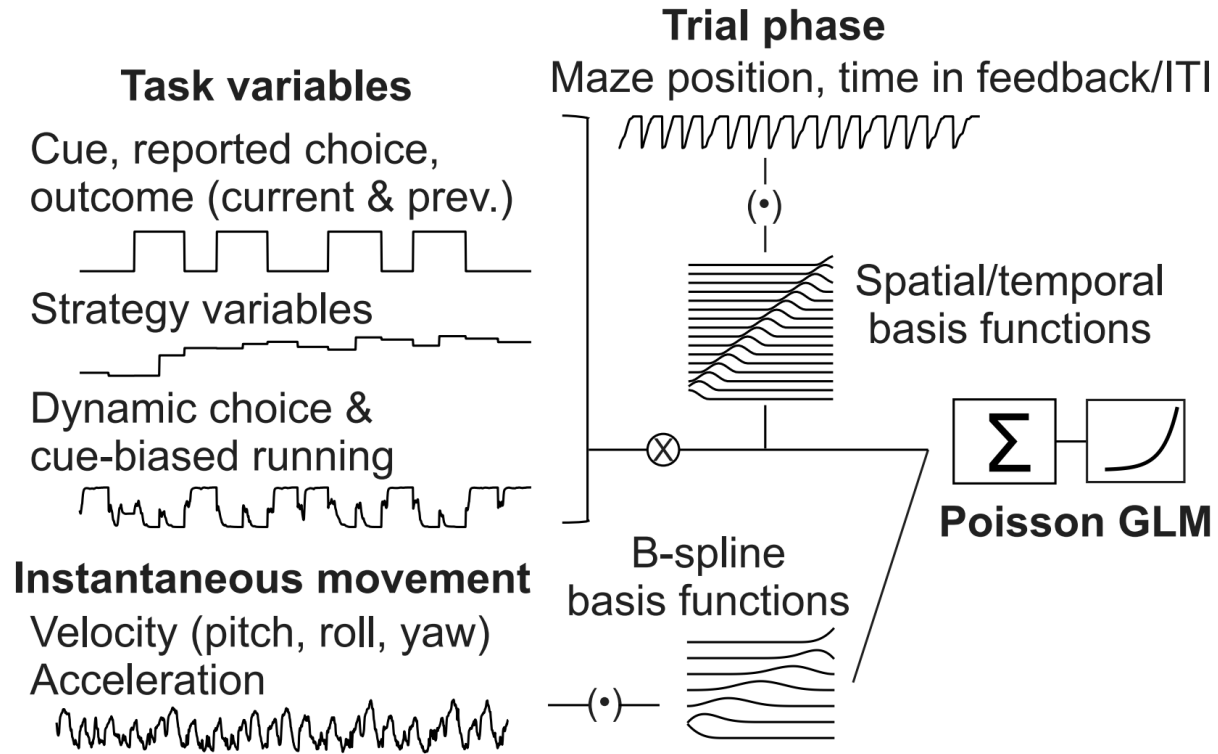
GLMs can also be used to characterize tuning in **dynamic tasks with nontrivial structure!**

See also tutorial slides by Shih-Yi at:  
[https://github.com/sytseng/GLM\\_Tensorflow\\_2](https://github.com/sytseng/GLM_Tensorflow_2)

Tseng and Chettih et al., *Neuron* (2022)



## Example 2: Fitting data from mouse posterior cortex



**Much more complex** for a few reasons:

### 1. Multiple categories of variables are relevant.

Dynamic navigation involves tracking:  
Movement (e.g., velocity), time (e.g., task phase),  
belief (go left or right?), strategy (e.g., be lazy?) ...

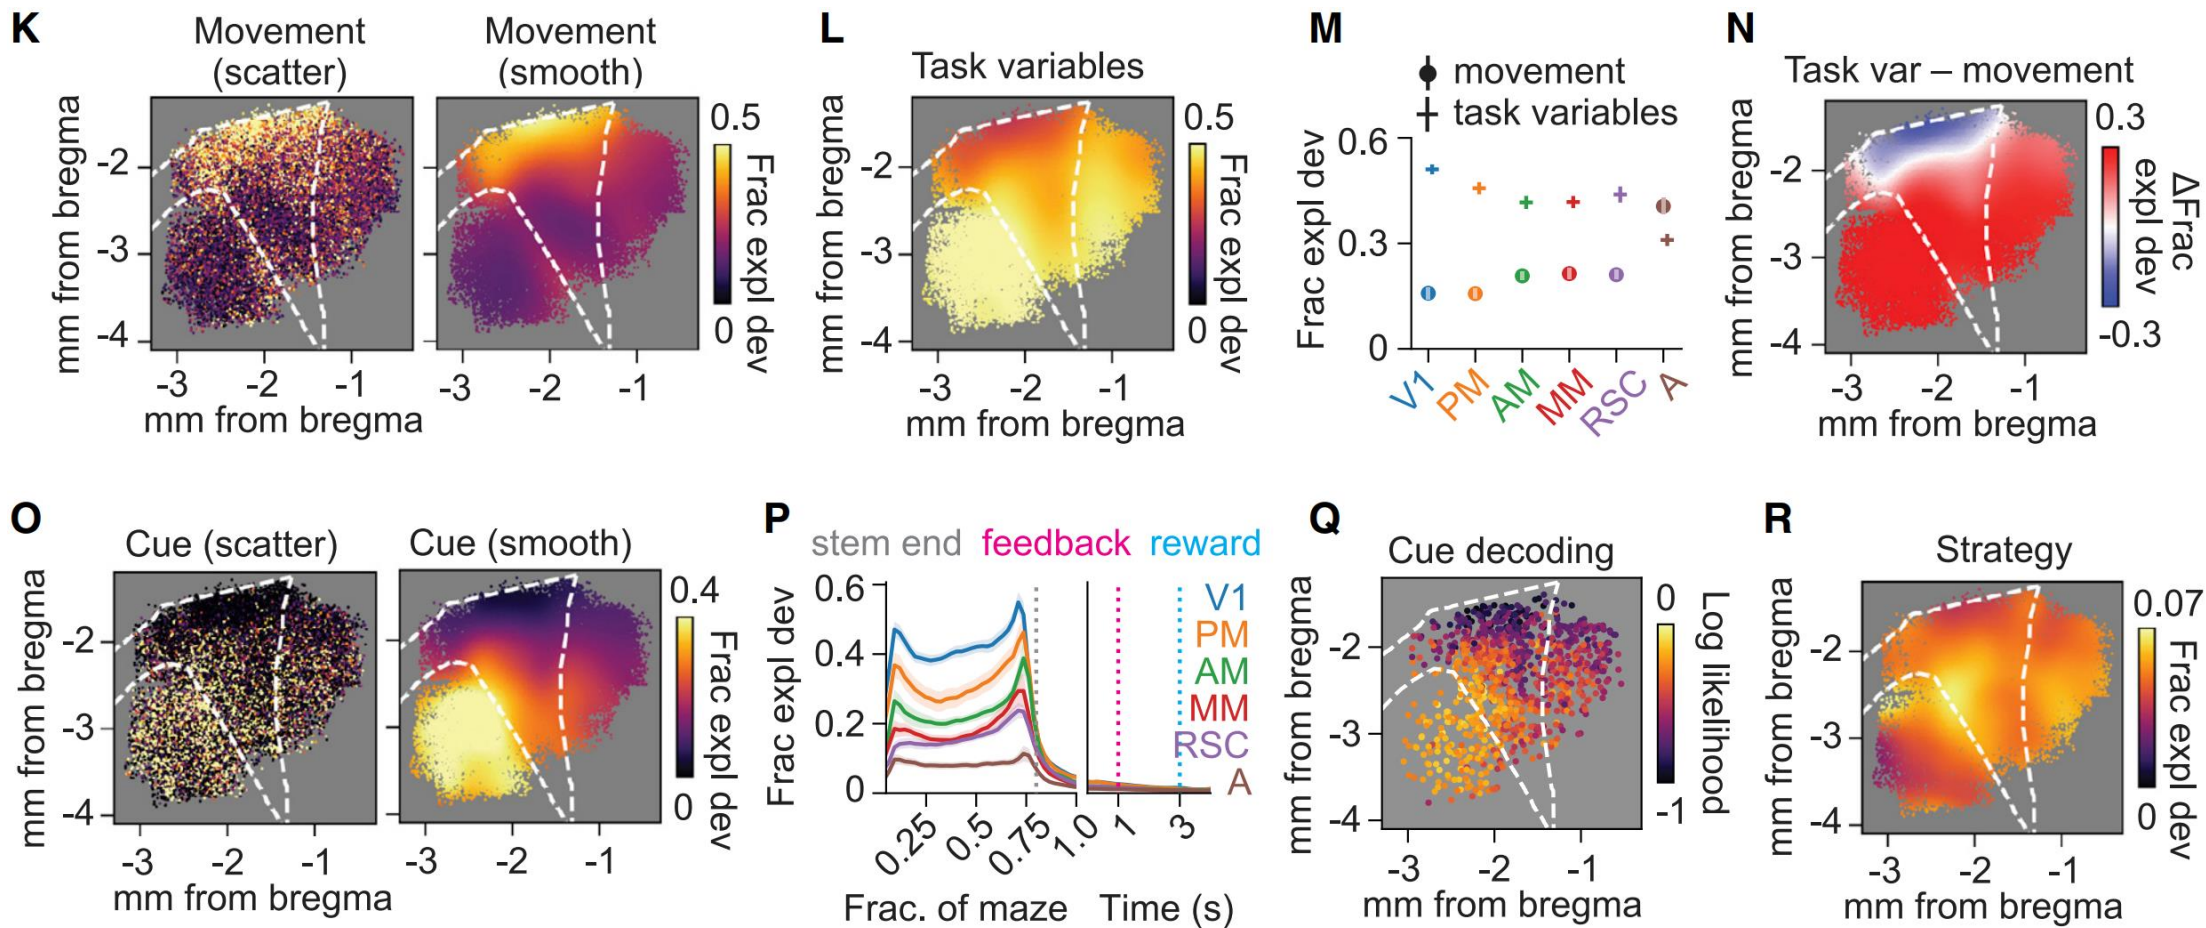
### 2. Time-dependence of firing rates is important.

e.g., evidence accumulation occurs over time,  
address using temporal basis functions

### 3. Parameter smoothness via state representation.

address using smooth basis functions instead of via regularization

## Example 2: Fitting data from mouse posterior cortex



## Example 2: Fitting data from mouse posterior cortex

### Important fitting detail:

Number of parameters is HUGE ( $\sim > 1000$  per neuron), so model highly underconstrained.

Tseng et al. used a group lasso penalty to regularize the model:

$$\lambda \sum_i \sqrt{g_i} \|\mathbf{w}_i\|_2$$

**Idea:** “encourages sparsity between tuning to different variables,  
but non-sparse L2 regularization on the within-variable bases”

Can also use some  
combination of:

$$\sum_j \frac{w_j^2}{2}$$

encourages  
small magnitudes

$$\sum_j |w_j|$$

encourages  
sparsity

## Example 2: Fitting data from mouse posterior cortex

Q. How do we measure goodness-of-fit?

For Poisson regression, better to use “Poisson deviance” than squared error.

$$D_{model} = 2 \sum_n \left[ y_n (\mathbf{w}^T \mathbf{x}_n + w_0 - \log \langle y \rangle) - (e^{\mathbf{w}^T \mathbf{x}_n + w_0} - \langle y \rangle) \right]$$

Can also consider a quantity analogous to “variance explained”: *deviance* explained!

$$R_{GLM}^2 := \frac{D_{model}}{D_{tot}} \geq 0 \quad \text{(see lecture notes for more details)}$$

**Idea:** Compare fitted model to two extremes. One extreme: only fit mean of data. (null model)  
Other extreme: # parameters = # data points (saturated model)

## Example 2: Fitting data from mouse posterior cortex

Q. How do we choose hyperparameters (e.g., regularization strengths)?

*Can't* just check which parameter fits better on test set!  
(because then you're using the test set as a training set...).

*Instead*, can use a strategy like k-fold cross validation.  
Split the training data up into pieces, and test model fits on one piece against fits on other pieces.

## Example 2 Takeaways

- Complex state representations allow GLMs to capture complicated spiking dynamics in a task that mixes dynamic navigation and decision-making.

On the other hand, additional model complexity has predictable tradeoffs, re: interpretability, time required to fit model, etc.

- Very large number of parameters means careful model selection is crucial.

In practice, ought to use tricks like regularization, k-fold cross-validation, etc...


Be careful about measuring goodness-of-fit; deviance more appropriate than squared error!

See Shih-Yi's GitHub repo for a research-level implementation of GLMs  
[https://github.com/sytseng/GLM\\_Tensorflow\\_2](https://github.com/sytseng/GLM_Tensorflow_2)

# Coding exercise: your turn!

<https://github.com/DrugowitschLab/ML-from-scratch-seminar>

**Goal:** Train Poisson GLM to fit data from 4 mouse MEC neurons



Data from Mallory and Hardcastle et al., *Nat Comm* (2021)

[https://figshare.com/authors/Lisa\\_Giocomo/9864194](https://figshare.com/authors/Lisa_Giocomo/9864194)